

Timers, Timing and PWM

Introduction to ARM Cortex-M Microprocessors – Chapter 10

Muním Zabidi

Faculty of Electrical Engineering, Universiti Teknologi Malaysia

- 1** Why This Matters
- 2** Why Software Delays Fail
- 3** Timer Architecture
- 4** Alternate Functions
- 5** PWM
- 6** Design Challenge

- 1** Why This Matters
- 2 Why Software Delays Fail
- 3 Timer Architecture
- 4 Alternate Functions
- 5 PWM
- 6 Design Challenge

Every program so far has used `delay_ms()` — a loop that counts down while doing nothing. This approach has three problems. Its timing can change with compiler optimization or clock frequency. It uses the CPU the whole time it waits. And it cannot give an accurate time base for measuring events. A hardware timer solves all three problems. It separates **timekeeping from program execution**. Once configured, it counts on its own, independent of the processor. By itself, it can **schedule** an event at a precise time, **measure** when an external event happens, or **generate** a precisely timed output — such as the PWM signal used to dim an LED or drive a motor. This chapter covers those three uses, in that order.

After completing this chapter, you should be able to:

1. Explain how a prescaler and auto-reload register determine a timer's period.
2. Calculate the prescaler and reload values required for a specified timing interval.
3. Configure a general-purpose timer to generate periodic update interrupts.
4. Use SysTick to maintain a millisecond system time base.
5. Select the correct alternate-function number to route a pin to a timer channel.
6. Calculate PWM frequency and duty cycle, and configure a timer output-compare channel to generate them.
7. Apply PWM to control LED brightness and motor speed.

- 1 Why This Matters
- 2 Why Software Delays Fail
- 3 Timer Architecture
- 4 Alternate Functions
- 5 PWM
- 6 Design Challenge

Three Things Wrong With `delay_ms()`


```
void delay_ms(int n) {  
    for (volatile int i = 0;  
        i < n * 4000; i++) ;  
}
```

1. Its timing changes with the **compiler optimisation level** and the **clock frequency**
2. It consumes **100%** of the CPU while waiting
3. It cannot be **interrupted accurately** – any ISR that fires makes the delay longer

A timer is a counter in hardware

It runs independently of the processor, at a rate you configure. It does not care what the CPU is doing, what the optimiser did, or how many interrupts fired. **That is what makes it a clock rather than a guess.**

- 1 Why This Matters
- 2 Why Software Delays Fail
- 3 **Timer Architecture**
- 4 Alternate Functions
- 5 PWM
- 6 Design Challenge



$$f_{\text{event}} = \frac{f_{\text{CK}}}{(\text{PSC} + 1) \times (\text{ARR} + 1)}$$

The **+1** is because a value of 0 means “divide by 1” – a counter that reloads at 0 still counts one tick.

Required

A 1 Hz interrupt from a 16 MHz clock.

Choose a prescaler

PSC = 15999 divides 16 MHz down to **1 kHz** — a convenient 1 ms tick.

Then the reload value

ARR = 999 gives 1000 ticks = **1 s**.

$$\text{Check: } \frac{16,000,000}{16000 \times 1000} = 1 \text{ Hz.}$$

TIM2->PSC = 15999;

TIM2->ARR = 999;

Why choose PSC first

Both registers are 16-bit, so neither can exceed 65535. Pick the prescaler to give a **convenient tick** — 1 ms or 1 μ s — then the reload value is just the number of ticks you want.

```
RCC->APB1ENR |= (1U << 0);    // TIM2 clock
TIM2->PSC = 15999;             // 16 MHz -> 1 kHz
TIM2->ARR = 999;               // 1000 ticks -> 1 Hz
TIM2->DIER |= (1U << 0);       // update interrupt enable
TIM2->CR1 |= (1U << 0);        // start counting
NVIC_EnableIRQ(TIM2_IRQn);

void TIM2_IRQHandler(void)
{
    if (TIM2->SR & (1U << 0)) { // update flag?
        TIM2->SR &= ~(1U << 0); // clear it
        GPIOA->ODR ^= (1U << 5);
    }
}
```

The CPU is now **free** between interrupts — it can sleep, or do real work, and the LED still blinks at exactly 1 Hz.

SysTick The system time base

- A **24-bit** down-counter built into the **core**, not the vendor's peripherals.
- Present on **every** Cortex-M, at the same address — so code using it is portable.
- Normally configured for a **1ms** tick.

```
SysTick_Config(SystemCoreClock / 1000);
```

```
volatile uint32_t ticks = 0;

void SysTick_Handler(void)
{
    ticks++;
}

uint32_t millis(void)
{
    return ticks;
}
```

This is what replaces `delay_ms()`

Instead of *waiting*, record when something last happened and check whether enough time has passed. That is the **non-blocking** pattern Chapter 14 builds a whole system from.

The Bad `delay_ms()` — For Comparison

```
void delay_ms_bad(int n)
{
    for (volatile int i = 0; i < n * 4000; i++)
        ;
    /* Wrong in three ways: drifts with optimisation level
       and clock frequency, burns 100% CPU, and any ISR that
       fires makes it longer than requested. Do not use this
       in real designs. */
}
```

A Proper `delay_ms()` — Built on SysTick

```
static volatile uint32_t systick_ms = 0;

void SysTick_Handler(void) { systick_ms++; }

void systick_init(uint32_t core_clock_hz)
{
    SysTick_Config(core_clock_hz / 1000U); // 1 ms tick
}

uint32_t millis(void) { return systick_ms; }

// Blocking, but exact regardless of compiler optimisation:
// it counts real elapsed ticks, not a fixed loop count.
void delay_ms(uint32_t n)
{
    uint32_t start = millis();
    while ((millis() - start) < n) {
        __asm volatile ("wfi"); // sleep between ticks
    }
}
```

An Alternative: `delay_ms()` on a General-Purpose Timer



```
void tim2_delay_init(void)
{
    RCC->APB1ENR |= (1U << 0); // TIM2 clock enable
    TIM2->PSC = 15999;          // 16 MHz -> 1 kHz
    TIM2->ARR = 0xFFFF;         // free-running
    TIM2->CR1 |= (1U << 0);      // start counting
}

void delay_ms_tim2(uint32_t n)
{
    TIM2->CNT = 0;
    while (TIM2->CNT < n)
        ;
}
```

Useful when SysTick is already claimed by an RTOS, and the mechanism is the same prescaler/auto-reload arithmetic from earlier in this chapter.

- 1 Why This Matters
- 2 Why Software Delays Fail
- 3 Timer Architecture
- 4 Alternate Functions**
- 5 PWM
- 6 Design Challenge

Routing a Pin to a Peripheral



A pin can serve several peripherals. MODER set to 10 selects alternate-function mode; AFR then chooses which of up to sixteen functions.

Selecting the Function Number

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16
AFRL7[3:0]				AFRL6[3:0]				AFRL5[3:0]				AFRL4[3:0]			
rw	rw	rw	rw	rw	rw	rw	rw	rw	rw	rw	rw	rw	rw	rw	rw
15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
AFRL3[3:0]				AFRL2[3:0]				AFRL1[3:0]				AFRL0[3:0]			
rw	rw	rw	rw	rw	rw	rw	rw	rw	rw	rw	rw	rw	rw	rw	rw

Four bits per pin, so AFR[0] covers pins 0–7 and AFR[1] covers pins 8–15.

For pin $n \geq 8$ the shift is $(n - 8) \times 4$.


```
// PA5 -> TIM2_CH1 is AF1
GPIOA->MODER &= ~(3U << 10);
GPIOA->MODER |= (2U << 10);

GPIOA->AFR[0] |= (1U << (5*4));
```

The function numbers are in the **datasheet**, not the reference manual – a distinction that costs people an afternoon.

- 1 Why This Matters
- 2 Why Software Delays Fail
- 3 Timer Architecture
- 4 Alternate Functions
- 5 PWM**
- 6 Design Challenge

Pulse Width Modulation



$$\text{duty cycle} = \frac{\text{on-time}}{T} \quad f_{\text{PWM}} = \frac{f_{\text{CK}}}{(\text{PSC} + 1)(\text{ARR} + 1)}$$

A digital pin producing an analog effect

The pin is only ever fully on or fully off — but switch it fast enough and an LED looks **dimmed**, a motor runs **slower**, a heater delivers **less power**. The average is what the load responds to.

Configuring PWM

```
TIM2->PSC  = 15;           // 16 MHz -> 1 MHz
TIM2->ARR   = 999;          // 1 MHz / 1000 = 1 kHz PWM
TIM2->CCR1  = 250;          // 250/1000 = 25% duty

TIM2->CCMR1 |= (6U << 4);   // PWM mode 1 on channel 1
TIM2->CCMR1 |= (1U << 3);   // preload enable
TIM2->CCER  |= (1U << 0);   // enable the output
TIM2->CR1   |= (1U << 0);   // start
```

ARR sets the frequency, CCR sets the duty

Change CCR1 at any time and the duty cycle changes on the next cycle — no glitch, because the preload register updates only at the reload point. **That is what makes smooth dimming possible.**

A higher PWM frequency is less visible but wastes more switching energy. For LEDs, 1 kHz is comfortably above flicker perception.

```
void pwm_init_tim2_ch1(uint32_t psc, uint32_t arr)
{
    RCC->APB1ENR |= (1U << 0);
    TIM2->PSC = psc;
    TIM2->ARR = arr;
    TIM2->CCR1 = 0;           // start at 0% duty
    TIM2->CCMR1 |= (6U << 4); // PWM mode 1, channel 1
    TIM2->CCMR1 |= (1U << 3); // preload enable
    TIM2->CCER |= (1U << 0);  // enable output
    TIM2->CR1  |= (1U << 0);  // start
}


// percent: 0-100
void pwm_set_duty_percent(uint8_t percent)
{
    if (percent > 100) percent = 100;
    uint32_t arr = TIM2->ARR;
    TIM2->CCR1 = ((uint32_t)percent * (arr + 1)) / 100U;
}
```

Example: A Breathing LED, Non-Blocking

```
void led_breathe_demo(void)
{
    static periodic_task_t step = { .last_ms = 0, .period_ms = 20 };
    static int duty = 0;
    static int dir = 1;

    if (task_due(&step)) {
        duty += dir;
        if (duty >= 100 || duty <= 0) dir = -dir;
        pwm_set_duty_percent((uint8_t)duty);
    }
}
```

Every 20 ms the duty cycle steps by one percentage point up or down, driven entirely by `task_due()` — no blocking call anywhere in the loop.



- A software delay is wrong three ways: it depends on the **optimiser** and **clock**, it burns the **whole CPU**, and interrupts **stretch** it.
- A timer is a **hardware counter**: $f = f_{\text{CK}} / ((\text{PSC} + 1)(\text{ARR} + 1))$.
- Choose **PSC** for a convenient tick, then **ARR** is the number of ticks.
- **SysTick** is in the core, identical on every Cortex-M, and is what a portable `millis()` is built on.
- An **alternate function** needs `MODER = 10` and the right number in `AFR` — four bits per pin.
- **PWM**: `ARR` sets the frequency, `CCR` sets the duty cycle. A digital pin produces an analog effect.

Next: Chapter 11 — ADC and Analog Interfacing

You can produce an analog effect. Next: measuring a real one.

- 1 Why This Matters
- 2 Why Software Delays Fail
- 3 Timer Architecture
- 4 Alternate Functions
- 5 PWM
- 6 Design Challenge

Design Challenge To be worked on paper, before any code

The brief

A system must perform three periodic tasks: read a sensor every 10 ms, update a display every 100 ms, and write a log entry every 1 s.

Design the timing structure using a single hardware time base. Specify the timer's PSC and ARR values for a 16 MHz clock, and show how one interrupt can schedule all three tasks without drift. Then work out what happens if the log write takes 40 ms, identify which requirement breaks first, and give two different ways to fix it — one that changes the schedule and one that changes the architecture.

Extension — a listening device. Now add a fourth activity to the same processor: sampling a microphone at 16 kHz for keyword spotting, as in audio timing. On paper, decide whether the system above can still meet all its deadlines. Produce a cycle budget at 16 MHz showing what fraction of the processor the sampling ISR alone uses. Decide which of the four activities may be delayed and which may not, and justify the ranking in terms of what breaks when each one slips. State where the audio buffer lives and how the main loop knows a window is ready without racing the ISR. Finally, identify the one activity in the original three that is now unsafe to leave as written, and say what you would change about it, without adding a second processor.

Design-Challenge Skeleton: One Time Base, Three Tasks

```
static periodic_task_t sensor_task = { .last_ms=0, .period_ms=10  };
static periodic_task_t display_task = { .last_ms=0, .period_ms=100 };
static periodic_task_t log_task    = { .last_ms=0, .period_ms=1000};

void read_sensor(void)    { /* ... */ }
void update_display(void) { /* ... */ }
void write_log(void)      { /* ... */ }

void scheduler_loop(void)
{
    for (;;) {
        if (task_due(&sensor_task)) read_sensor();
        if (task_due(&display_task)) update_display();
        if (task_due(&log_task))    write_log();
    }
}
```

One SysTick-driven `millis()` schedules the sensor (10 ms), display (100 ms) and log (1 s) tasks from a single interrupt source, without drift.

Questions?

Introduction to ARM Cortex-M Microprocessors — Chapter 10